

Churn_Analysis_

November 12, 2022

1 Churn Customers

1.0.1 Introduction

This IBM Sample Dataset has information about Telco customers and if they left the company within the last month (churn).

Basic information: Customers who left within the last month – the column is called Churn. Services that each customer has signed up for – phone, multiple lines, internet, online security, online backup, device protection, tech support, and streaming TV and movies. Customer account information – how long they’ve been a customer, contract, payment method, paperless billing, monthly charges, and total charges. Demographic info about customers – gender, age range, and if they have partners and dependents. There are 21 columns with 19 features.

Objective I will explore the data and try to answer some questions like:

Customer churn measures how and why are customers leaving the business

1.0.2 Importing Libraries

```
[ ]: import pandas as pd
      from matplotlib import pyplot as plt
      import seaborn as sns
      import tensorflow as tf
      from tensorflow import keras
```

1.0.3 Loading data

```
[ ]: df = pd.read_csv("Churndata.csv")
      df.sample(5)
```

```
[ ]:      customerID  gender  SeniorCitizen  ...  MonthlyCharges  TotalCharges  Churn
      3821  1833-VGRUM  Female                1  ...           74.20           74.2    Yes
      4831  4654-GGUII  Female                0  ...           40.20          711.95     No
      6528  4957-SREEC   Male                0  ...           20.35          1458.1     No
      6126  9190-MFJLN   Male                1  ...           95.90          1777.9     Yes
      4267  1227-UDMZR  Female                0  ...           91.15          6637.9     No
```

[5 rows x 21 columns]

1.0.4 Pre-processing

```
[ ]: #Checking for datatypes
#We see that Total charges is an object. We need to change that to float
df.dtypes
```

```
[ ]: customerID      object
gender             object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
OnlineBackup       object
DeviceProtection   object
TechSupport        object
StreamingTV        object
StreamingMovies    object
Contract           object
PaperlessBilling   object
PaymentMethod      object
MonthlyCharges     float64
TotalCharges       object
Churn              object
dtype: object
```

```
[ ]: df.TotalCharges.values
```

```
[ ]: array(['29.85', '1889.5', '108.15', ..., '346.45', '306.6', '6844.5'],
      dtype=object)
```

```
[ ]: df['TotalCharges'] = df['TotalCharges'].replace(" ", 0).astype('float64')
```

```
[ ]: df.dtypes
```

```
[ ]: customerID      object
gender             object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
```

```

OnlineBackup      object
DeviceProtection  object
TechSupport       object
StreamingTV       object
StreamingMovies   object
Contract          object
PaperlessBilling  object
PaymentMethod     object
MonthlyCharges    float64
TotalCharges      float64
Churn             object
dtype: object

```

1.0.5 Data Visualization

Viz1 : No. of customers Vs. Tenure

```

[ ]: tenure_churn_no = df[df.Churn=='No'].tenure
    tenure_churn_yes = df[df.Churn=='Yes'].tenure

plt.xlabel("tenure")
plt.ylabel("Number Of Customers")
plt.title("Customer Churn Prediction Visualization")

blood_sugar_men = [113, 85, 90, 150, 149, 88, 93, 115, 135, 80, 77, 82, 129]
blood_sugar_women = [67, 98, 89, 120, 133, 150, 84, 69, 89, 79, 120, 112, 100]

plt.hist([tenure_churn_yes, tenure_churn_no], rwidth=0.95,
        color=['green', 'red'], label=['Churn=Yes', 'Churn=No'])
plt.legend()


#In this graph, we have the number of customers who will be leaving and not
↳leaving the company, and we have the tenure of months that they have been
↳apart of the company till now.

```

```

[ ]: <matplotlib.legend.Legend at 0x7f877bf79cf8>


```



Churn Vs. Count Churn: No - 72.4%


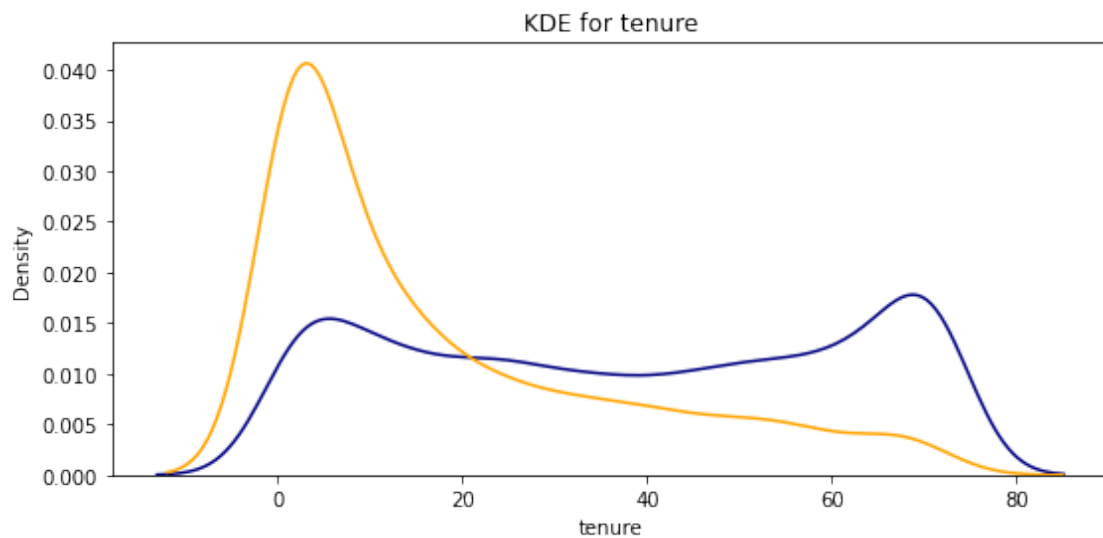
Churn: Yes - 27.6%

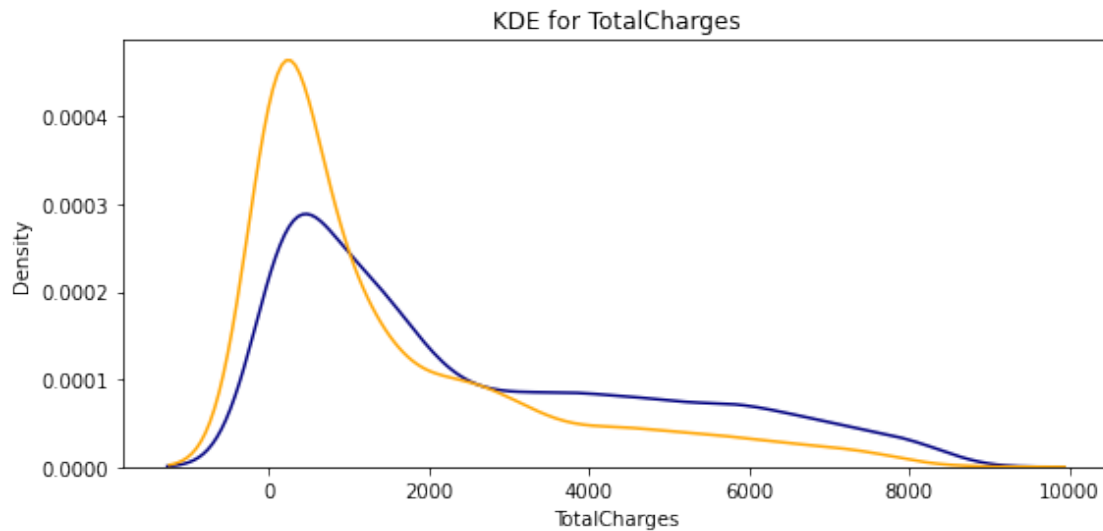
```
[ ]: ax = sns.catplot(y="Churn", kind="count", data=df, height=2.6, aspect=2.5,
    ↪orient='h')
```



Numeric Features There are only three numerical columns: tenure, monthly charges and total charges. The probability density distribution can be estimate using the seaborn kdeplot function.

```
[ ]: def kdeplot(feature):  
    plt.figure(figsize=(9, 4))  
    plt.title("KDE for {}".format(feature))  
    ax0 = sns.kdeplot(df[df['Churn'] == 'No'][feature].dropna(), color= 'navy',  
↳label= 'Churn: No')  
    ax1 = sns.kdeplot(df[df['Churn'] == 'Yes'][feature].dropna(), color=  
↳'orange', label= 'Churn: Yes')  
kdeplot('tenure')  
kdeplot('MonthlyCharges')  
kdeplot('TotalCharges')
```





From the plots above we can conclude that:

Recent clients are more likely to churn.

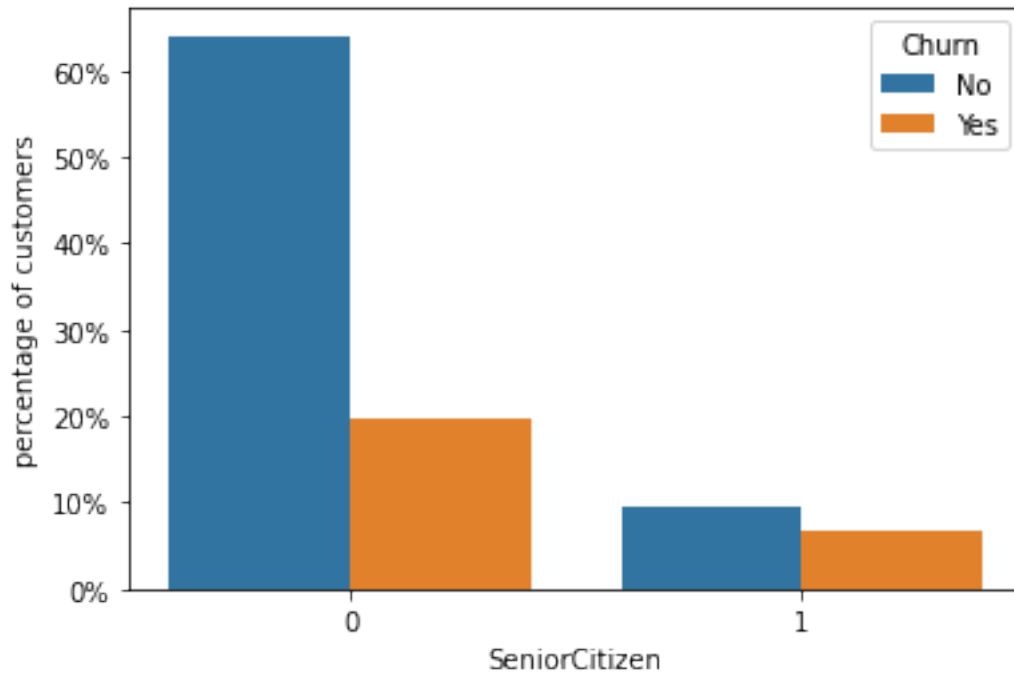
Clients with higher MonthlyCharges are also more likely to churn.

Tenure and MonthlyCharges are probably important features.


Categorical features This dataset has 16 categorical features.

Gender and Age(SeniorCitizen)

```
[ ]: def barplot_percentages(feature, orient='v', axis_name="percentage of_
    ↪customers"):
    ratios = pd.DataFrame()
    g = df.groupby(feature)["Churn"].value_counts().to_frame()
    g = g.rename({"Churn": axis_name}, axis=1).reset_index()
    g[axis_name] = g[axis_name]/len(df)
    if orient == 'v':
        ax = sns.barplot(x=feature, y= axis_name, hue='Churn', data=g,
    ↪orient=orient)
        ax.set_yticklabels(['{:.0%}'.format(y) for y in ax.get_yticks()])
    else:
        ax = sns.barplot(x= axis_name, y=feature, hue='Churn', data=g,
    ↪orient=orient)
        ax.set_xticklabels(['{:.0%}'.format(x) for x in ax.get_xticks()])
    ax.plot()
    barplot_percentages("SeniorCitizen")
```




```
[ ]: df['churn_rate'] = df['Churn'].replace("No", 0).replace("Yes", 1)
g = sns.FacetGrid(df, col="SeniorCitizen", height=4, aspect=.9)
ax = g.map(sns.barplot, "gender", "churn_rate", palette = "Blues_d", order=[
    'Female', 'Male'])
```



Gender is not an indicative of churn. SeniorCitizens are only 16% of customers, but they have a much higher churn rate: 42% against 23% for non-senior customers.

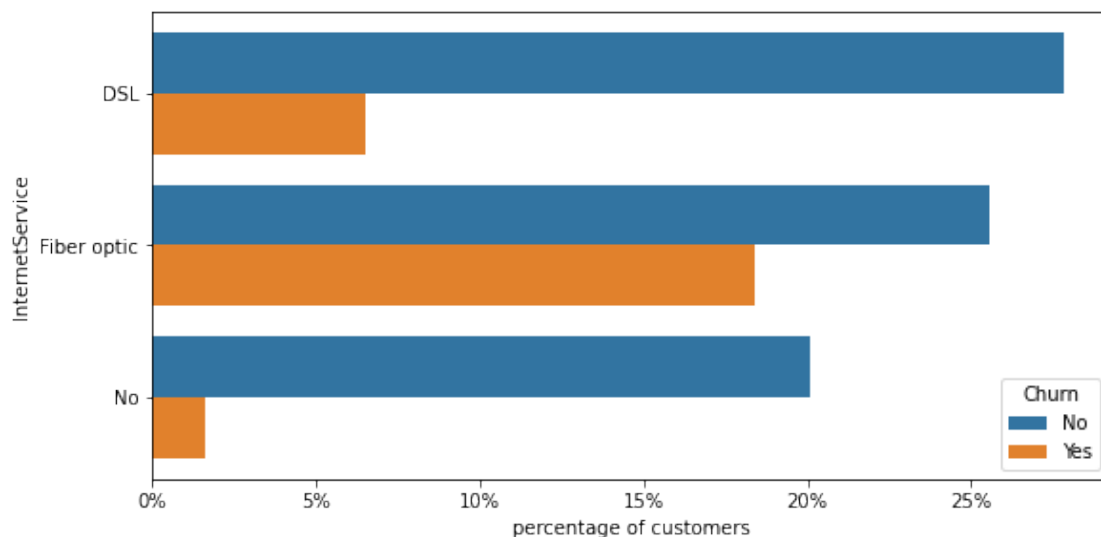
Phone and Internet services There are only two features here: If the client has phone and if he has more than one line. Both can be summed up in one chart:

```
[ ]: plt.figure(figsize=(9, 4.5))  
      barplot_percentages("MultipleLines", orient='h')
```




Customer with multiple lines have higher of churn rate.

```
[ ]: plt.figure(figsize=(9, 4.5))  
      barplot_percentages("InternetService", orient="h")
```




Additional service analysis The first plot shows the total number of customers for each additional service, while the second shows the number of clients that churn.

```
[ ]: #There are six additional services for customers with internet:
cols = ["OnlineSecurity", "OnlineBackup", "DeviceProtection", "TechSupport",
        "StreamingTV", "StreamingMovies"]
df1 = pd.melt(df[df["InternetService"] != "No"][cols]).rename({'value': 'Has_
        service'}, axis=1)
plt.figure(figsize=(10, 4.5))
ax = sns.countplot(data=df1, x='variable', hue='Has service')
ax.set(xlabel='Additional service', ylabel='Num of customers')
plt.show()
```



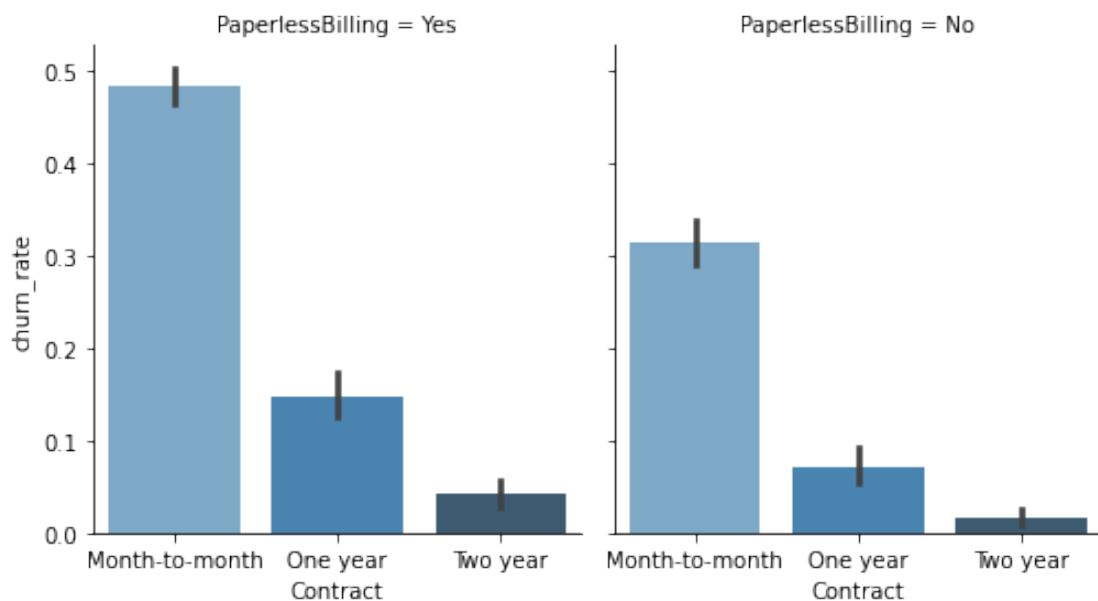
```
[ ]: plt.figure(figsize=(10, 4.5))
df1 = df[(df.InternetService != "No") & (df.Churn == "Yes")]
df1 = pd.melt(df1[cols]).rename({'value': 'Has service'}, axis=1)
ax = sns.countplot(data=df1, x='variable', hue='Has service', hue_order=['No',
        'Yes'])
ax.set(xlabel='Additional service', ylabel='Num of churns')
plt.show()
```




Customers with the OnlineSecurity, Backup, Protection and tech support are more unlikely to churn.

Contract and Payment Customers with paperless billing are more probable to churn. The preferred payment method is Electronic check with around 35% of customers. This also has a very high churn rate.

```
[ ]: g = sns.FacetGrid(df, col="PaperlessBilling", height=4, aspect=.9)
ax = g.map(sns.barplot, "Contract", "churn_rate", palette = "Blues_d", order=
↳ ['Month-to-month', 'One year', 'Two year'])
```




```
[ ]: plt.figure(figsize=(9, 4.5))
      barplot_percentages("PaymentMethod", orient='h')
```



Correlation between features

```
[ ]: plt.figure(figsize=(12, 6))
      df.drop(['customerID', 'churn_rate'],
              axis=1, inplace=True)
      corr = df.apply(lambda x: pd.factorize(x)[0]).corr()
      ax = sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns,
                      linewidths=.2, cmap="YlGnBu")
```



1.0.6 Pre-process to build model

```
[ ]: df.head()
```

```
[ ]:      gender  SeniorCitizen  Partner  ...  MonthlyCharges  TotalCharges  Churn
0  Female           0      Yes  ...      29.85         29.85    No
1   Male           0      No  ...      56.95        1889.50    No
2   Male           0      No  ...      53.85         108.15   Yes
3   Male           0      No  ...      42.30        1840.75    No
4  Female           0      No  ...      70.70         151.65   Yes
```

[5 rows x 20 columns]

```
[ ]: #Let's print unique values in object columns to see data values
def print_unique_col_values(df):
    for column in df:
        if df[column].dtypes=='object':
            print(f'{column}: {df[column].unique()}')
```

```
[ ]: print_unique_col_values(df)
```

```
gender: ['Female' 'Male']
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
PhoneService: ['No' 'Yes']
MultipleLines: ['No phone service' 'No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes' 'No internet service']
OnlineBackup: ['Yes' 'No' 'No internet service']
DeviceProtection: ['No' 'Yes' 'No internet service']
TechSupport: ['No' 'Yes' 'No internet service']
StreamingTV: ['No' 'Yes' 'No internet service']
StreamingMovies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
 'Credit card (automatic)']
Churn: ['No' 'Yes']
```

```
[ ]: #Some of the columns have no internet service or no phone service, that can be
      ↪replaced with a simple No
df.replace('No internet service','No',inplace=True)
df.replace('No phone service','No',inplace=True)
```

```
[ ]: print_unique_col_values(df)
```

```

gender: ['Female' 'Male']
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
PhoneService: ['No' 'Yes']
MultipleLines: ['No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes']
OnlineBackup: ['Yes' 'No']
DeviceProtection: ['No' 'Yes']
TechSupport: ['No' 'Yes']
StreamingTV: ['No' 'Yes']
StreamingMovies: ['No' 'Yes']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
Churn: ['No' 'Yes']

```

```
[ ]: df
```

```

[ ]:
   gender  SeniorCitizen  Partner  ...  MonthlyCharges  TotalCharges  Churn
0   Female              0     Yes  ...           29.85         29.85    No
1     Male              0     No   ...           56.95        1889.50    No
2     Male              0     No   ...           53.85         108.15   Yes
3     Male              0     No   ...           42.30        1840.75    No
4   Female              0     No   ...           70.70         151.65   Yes
...     ...            ...     ...  ...             ...           ...
7038  Male              0     Yes  ...           84.80        1990.50    No
7039  Female              0     Yes  ...          103.20        7362.90    No
7040  Female              0     Yes  ...           29.60         346.45    No
7041  Male              1     Yes  ...           74.40         306.60   Yes
7042  Male              0     No   ...          105.65        6844.50    No

```

[7043 rows x 20 columns]

```
[ ]: #Here, converting Yes= 1, No =0
```

```

yes_no_columns =
    ↪ ['Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'OnlineSecurity', 'OnlineBackup',
      ↪
    ↪ 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'PaperlessBilling', 'Churn']
for col in yes_no_columns:
    df[col].replace({'Yes': 1, 'No': 0}, inplace=True)

```

```

[ ]: for col in df:
      print(f'{col}: {df[col].unique()}')

```

```
gender: ['Female' 'Male']
```

```

SeniorCitizen: [0 1]
Partner: [1 0]
Dependents: [0 1]
tenure: [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27
  5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
 32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26  0
 39]
PhoneService: [0 1]
MultipleLines: [0 1]
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: [0 1]
OnlineBackup: [1 0]
DeviceProtection: [0 1]
TechSupport: [0 1]
StreamingTV: [0 1]
StreamingMovies: [0 1]
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: [1 0]
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
  'Credit card (automatic)']
MonthlyCharges: [29.85 56.95 53.85 ... 63.1  44.2  78.7 ]
TotalCharges: [ 29.85 1889.5  108.15 ...  346.45  306.6  6844.5 ]
Churn: [0 1]

```

```
[ ]: df['gender'].replace({'Female':1,'Male':0},inplace=True)
```

```
[ ]: df.gender.unique()
```

```
[ ]: array([1, 0])
```

One hot encoding for categorical columns

```
[ ]: df1 = pd.get_dummies(data=df,
    ↪columns=['InternetService','Contract','PaymentMethod'])
df1.columns
```

```
[ ]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
  'PhoneService', 'MultipleLines', 'OnlineSecurity', 'OnlineBackup',
  'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
  'PaperlessBilling', 'MonthlyCharges', 'TotalCharges', 'Churn',
  'InternetService_DSL', 'InternetService_Fiber optic',
  'InternetService_No', 'Contract_Month-to-month', 'Contract_One year',
  'Contract_Two year', 'PaymentMethod_Bank transfer (automatic)',
  'PaymentMethod_Credit card (automatic)',
  'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check'],
  dtype='object')
```

```
[ ]: df1.sample(5)
```

```
[ ]:      gender  ...  PaymentMethod_Mailed check
6383      0  ...                      1
5006      1  ...                      1
1189      0  ...                      0
1956      1  ...                      0
4435      0  ...                      0
```

[5 rows x 27 columns]

```
[ ]: df1.dtypes
```

```
[ ]: gender                int64
SeniorCitizen             int64
Partner                   int64
Dependents                 int64
tenure                     int64
PhoneService              int64
MultipleLines             int64
OnlineSecurity            int64
OnlineBackup              int64
DeviceProtection          int64
TechSupport               int64
StreamingTV               int64
StreamingMovies           int64
PaperlessBilling          int64
MonthlyCharges            float64
TotalCharges              float64
Churn                     int64
InternetService_DSL       uint8
InternetService_Fiber optic uint8
InternetService_No        uint8
Contract_Month-to-month   uint8
Contract_One year         uint8
Contract_Two year         uint8
PaymentMethod_Bank transfer (automatic) uint8
PaymentMethod_Credit card (automatic)   uint8
PaymentMethod_Electronic check          uint8
PaymentMethod_Mailed check              uint8
dtype: object
```

```
[ ]: cols_to_scale = ['tenure', 'MonthlyCharges', 'TotalCharges']

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df1[cols_to_scale] = scaler.fit_transform(df1[cols_to_scale])
```

```
[ ]: for col in df1:
      print(f'{col}: {df1[col].unique()}')
```

```
gender: [1 0]
SeniorCitizen: [0 1]
Partner: [1 0]
Dependents: [0 1]
tenure: [0.01388889 0.47222222 0.02777778 0.625          0.11111111 0.30555556
 0.13888889 0.38888889 0.86111111 0.18055556 0.22222222 0.80555556
 0.68055556 0.34722222 0.95833333 0.72222222 0.98611111 0.29166667
 0.16666667 0.41666667 0.65277778 1.          0.23611111 0.375
 0.06944444 0.63888889 0.15277778 0.97222222 0.875          0.59722222
 0.20833333 0.83333333 0.25          0.91666667 0.125          0.04166667
 0.43055556 0.69444444 0.88888889 0.77777778 0.09722222 0.58333333
 0.48611111 0.66666667 0.40277778 0.90277778 0.52777778 0.94444444
 0.44444444 0.76388889 0.51388889 0.5          0.56944444 0.08333333
 0.05555556 0.45833333 0.93055556 0.31944444 0.79166667 0.84722222
 0.19444444 0.27777778 0.73611111 0.55555556 0.81944444 0.33333333
 0.61111111 0.26388889 0.75          0.70833333 0.36111111 0.
 0.54166667]
PhoneService: [0 1]
MultipleLines: [0 1]
OnlineSecurity: [0 1]
OnlineBackup: [1 0]
DeviceProtection: [0 1]
TechSupport: [0 1]
StreamingTV: [0 1]
StreamingMovies: [0 1]
PaperlessBilling: [1 0]
MonthlyCharges: [0.11542289 0.38507463 0.35422886 ... 0.44626866 0.25820896
0.60149254]
TotalCharges: [0.00343704 0.21756402 0.01245279 ... 0.03989153 0.03530306
0.78810105]
Churn: [0 1]
InternetService_DSL: [1 0]
InternetService_Fiber optic: [0 1]
InternetService_No: [0 1]
Contract_Month-to-month: [1 0]
Contract_One year: [0 1]
Contract_Two year: [0 1]
PaymentMethod_Bank transfer (automatic): [0 1]
PaymentMethod_Credit card (automatic): [0 1]
PaymentMethod_Electronic check: [1 0]
PaymentMethod_Mailed check: [0 1]
```

Train test split


```
[ ]: X = df1.drop('Churn',axis='columns')
y = df1['Churn']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↪2,random_state=5)
```

```
[ ]: X_train.shape
```

```
[ ]: (5634, 26)
```

```
[ ]: X_test.shape
```

```
[ ]: (1409, 26)
```

```
[ ]: X_train[:10]
```

```
[ ]:      gender  ...  PaymentMethod_Mailed check
5860      1  ...
2458      0  ...
5879      0  ...
4708      1  ...
1293      0  ...
2242      0  ...
1444      0  ...
3269      0  ...
101       1  ...
4191      1  ...
```

```
[10 rows x 26 columns]
```

```
[ ]: len(X_train.columns)
```

```
[ ]: 26
```

1.0.7 Build a model (ANN) in tensorflow/keras

```
[ ]: model = keras.Sequential([
    keras.layers.Dense(26, input_shape=(26,), activation='relu'),
    keras.layers.Dense(15, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])

# opt = keras.optimizers.Adam(learning_rate=0.01)

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
model.fit(X_train, y_train, epochs=100)
```

```
Epoch 1/100
177/177 [=====] - 0s 1ms/step - loss: 0.5035 -
accuracy: 0.7474
Epoch 2/100
177/177 [=====] - 0s 1ms/step - loss: 0.4302 -
accuracy: 0.7930
Epoch 3/100
177/177 [=====] - 0s 1ms/step - loss: 0.4229 -
accuracy: 0.7980
Epoch 4/100
177/177 [=====] - 0s 1ms/step - loss: 0.4181 -
accuracy: 0.8055
Epoch 5/100
177/177 [=====] - 0s 1ms/step - loss: 0.4145 -
accuracy: 0.8067
Epoch 6/100
177/177 [=====] - 0s 1ms/step - loss: 0.4132 -
accuracy: 0.8046
Epoch 7/100
177/177 [=====] - 0s 1ms/step - loss: 0.4118 -
accuracy: 0.8090
Epoch 8/100
177/177 [=====] - 0s 1ms/step - loss: 0.4093 -
accuracy: 0.8080
Epoch 9/100
177/177 [=====] - 0s 1ms/step - loss: 0.4086 -
accuracy: 0.8094
Epoch 10/100
177/177 [=====] - 0s 1ms/step - loss: 0.4080 -
accuracy: 0.8083
Epoch 11/100
177/177 [=====] - 0s 1ms/step - loss: 0.4061 -
accuracy: 0.8080
Epoch 12/100
177/177 [=====] - 0s 1ms/step - loss: 0.4037 -
accuracy: 0.8126
Epoch 13/100
177/177 [=====] - 0s 1ms/step - loss: 0.4040 -
accuracy: 0.8131
Epoch 14/100
177/177 [=====] - 0s 1ms/step - loss: 0.4034 -
accuracy: 0.8115
Epoch 15/100
177/177 [=====] - 0s 1ms/step - loss: 0.4013 -
accuracy: 0.8131
```

Epoch 16/100
177/177 [=====] - 0s 1ms/step - loss: 0.4001 -
accuracy: 0.8168
Epoch 17/100
177/177 [=====] - 0s 1ms/step - loss: 0.4010 -
accuracy: 0.8131
Epoch 18/100
177/177 [=====] - 0s 1ms/step - loss: 0.3993 -
accuracy: 0.8154
Epoch 19/100
177/177 [=====] - 0s 1ms/step - loss: 0.3996 -
accuracy: 0.8156
Epoch 20/100
177/177 [=====] - 0s 1ms/step - loss: 0.3985 -
accuracy: 0.8165
Epoch 21/100
177/177 [=====] - 0s 1ms/step - loss: 0.3971 -
accuracy: 0.8175
Epoch 22/100
177/177 [=====] - 0s 1ms/step - loss: 0.3961 -
accuracy: 0.8181
Epoch 23/100
177/177 [=====] - 0s 1ms/step - loss: 0.3952 -
accuracy: 0.8154
Epoch 24/100
177/177 [=====] - 0s 1ms/step - loss: 0.3941 -
accuracy: 0.8175
Epoch 25/100
177/177 [=====] - 0s 1ms/step - loss: 0.3955 -
accuracy: 0.8175
Epoch 26/100
177/177 [=====] - 0s 1ms/step - loss: 0.3925 -
accuracy: 0.8197
Epoch 27/100
177/177 [=====] - 0s 1ms/step - loss: 0.3931 -
accuracy: 0.8170
Epoch 28/100
177/177 [=====] - 0s 1ms/step - loss: 0.3924 -
accuracy: 0.8179
Epoch 29/100
177/177 [=====] - 0s 1ms/step - loss: 0.3910 -
accuracy: 0.8172
Epoch 30/100
177/177 [=====] - 0s 1ms/step - loss: 0.3905 -
accuracy: 0.8193
Epoch 31/100
177/177 [=====] - 0s 1ms/step - loss: 0.3909 -
accuracy: 0.8206

Epoch 32/100
177/177 [=====] - 0s 1ms/step - loss: 0.3895 -
accuracy: 0.8195
Epoch 33/100
177/177 [=====] - 0s 1ms/step - loss: 0.3878 -
accuracy: 0.8181
Epoch 34/100
177/177 [=====] - 0s 1ms/step - loss: 0.3887 -
accuracy: 0.8179
Epoch 35/100
177/177 [=====] - 0s 1ms/step - loss: 0.3874 -
accuracy: 0.8202
Epoch 36/100
177/177 [=====] - 0s 1ms/step - loss: 0.3864 -
accuracy: 0.8191
Epoch 37/100
177/177 [=====] - 0s 1ms/step - loss: 0.3863 -
accuracy: 0.8200
Epoch 38/100
177/177 [=====] - 0s 1ms/step - loss: 0.3862 -
accuracy: 0.8181
Epoch 39/100
177/177 [=====] - 0s 1ms/step - loss: 0.3860 -
accuracy: 0.8179
Epoch 40/100
177/177 [=====] - 0s 1ms/step - loss: 0.3858 -
accuracy: 0.8211
Epoch 41/100
177/177 [=====] - 0s 1ms/step - loss: 0.3836 -
accuracy: 0.8200
Epoch 42/100
177/177 [=====] - 0s 1ms/step - loss: 0.3842 -
accuracy: 0.8223
Epoch 43/100
177/177 [=====] - 0s 1ms/step - loss: 0.3827 -
accuracy: 0.8204
Epoch 44/100
177/177 [=====] - 0s 1ms/step - loss: 0.3831 -
accuracy: 0.8220
Epoch 45/100
177/177 [=====] - 0s 1ms/step - loss: 0.3831 -
accuracy: 0.8222
Epoch 46/100
177/177 [=====] - 0s 1ms/step - loss: 0.3817 -
accuracy: 0.8198
Epoch 47/100
177/177 [=====] - 0s 1ms/step - loss: 0.3817 -
accuracy: 0.8214

Epoch 48/100
177/177 [=====] - 0s 1ms/step - loss: 0.3812 -
accuracy: 0.8222
Epoch 49/100
177/177 [=====] - 0s 1ms/step - loss: 0.3809 -
accuracy: 0.8227
Epoch 50/100
177/177 [=====] - 0s 1ms/step - loss: 0.3799 -
accuracy: 0.8211
Epoch 51/100
177/177 [=====] - 0s 1ms/step - loss: 0.3797 -
accuracy: 0.8218
Epoch 52/100
177/177 [=====] - 0s 1ms/step - loss: 0.3796 -
accuracy: 0.8220
Epoch 53/100
177/177 [=====] - 0s 1ms/step - loss: 0.3788 -
accuracy: 0.8232
Epoch 54/100
177/177 [=====] - 0s 1ms/step - loss: 0.3785 -
accuracy: 0.8246
Epoch 55/100
177/177 [=====] - 0s 1ms/step - loss: 0.3776 -
accuracy: 0.8211
Epoch 56/100
177/177 [=====] - 0s 1ms/step - loss: 0.3783 -
accuracy: 0.8236
Epoch 57/100
177/177 [=====] - 0s 1ms/step - loss: 0.3771 -
accuracy: 0.8245
Epoch 58/100
177/177 [=====] - 0s 1ms/step - loss: 0.3762 -
accuracy: 0.8220
Epoch 59/100
177/177 [=====] - 0s 1ms/step - loss: 0.3757 -
accuracy: 0.8248
Epoch 60/100
177/177 [=====] - 0s 1ms/step - loss: 0.3759 -
accuracy: 0.8237
Epoch 61/100
177/177 [=====] - 0s 1ms/step - loss: 0.3750 -
accuracy: 0.8268
Epoch 62/100
177/177 [=====] - 0s 1ms/step - loss: 0.3752 -
accuracy: 0.8243
Epoch 63/100
177/177 [=====] - 0s 1ms/step - loss: 0.3758 -
accuracy: 0.8255

Epoch 64/100
177/177 [=====] - 0s 1ms/step - loss: 0.3735 -
accuracy: 0.8250

Epoch 65/100
177/177 [=====] - 0s 1ms/step - loss: 0.3732 -
accuracy: 0.8271

Epoch 66/100
177/177 [=====] - 0s 1ms/step - loss: 0.3738 -
accuracy: 0.8248

Epoch 67/100
177/177 [=====] - 0s 1ms/step - loss: 0.3735 -
accuracy: 0.8229

Epoch 68/100
177/177 [=====] - 0s 1ms/step - loss: 0.3715 -
accuracy: 0.8236

Epoch 69/100
177/177 [=====] - 0s 1ms/step - loss: 0.3713 -
accuracy: 0.8241

Epoch 70/100
177/177 [=====] - 0s 1ms/step - loss: 0.3718 -
accuracy: 0.8252

Epoch 71/100
177/177 [=====] - 0s 1ms/step - loss: 0.3720 -
accuracy: 0.8259

Epoch 72/100
177/177 [=====] - 0s 1ms/step - loss: 0.3714 -
accuracy: 0.8241

Epoch 73/100
177/177 [=====] - 0s 1ms/step - loss: 0.3712 -
accuracy: 0.8264

Epoch 74/100
177/177 [=====] - 0s 1ms/step - loss: 0.3703 -
accuracy: 0.8262

Epoch 75/100
177/177 [=====] - 0s 1ms/step - loss: 0.3696 -
accuracy: 0.8230

Epoch 76/100
177/177 [=====] - 0s 1ms/step - loss: 0.3690 -
accuracy: 0.8282

Epoch 77/100
177/177 [=====] - 0s 1ms/step - loss: 0.3691 -
accuracy: 0.8285

Epoch 78/100
177/177 [=====] - 0s 1ms/step - loss: 0.3691 -
accuracy: 0.8252

Epoch 79/100
177/177 [=====] - 0s 1ms/step - loss: 0.3690 -
accuracy: 0.8259

Epoch 80/100
177/177 [=====] - 0s 1ms/step - loss: 0.3676 -
accuracy: 0.8259

Epoch 81/100
177/177 [=====] - 0s 1ms/step - loss: 0.3681 -
accuracy: 0.8239

Epoch 82/100
177/177 [=====] - 0s 1ms/step - loss: 0.3675 -
accuracy: 0.8296

Epoch 83/100
177/177 [=====] - 0s 1ms/step - loss: 0.3661 -
accuracy: 0.8236

Epoch 84/100
177/177 [=====] - 0s 1ms/step - loss: 0.3651 -
accuracy: 0.8278

Epoch 85/100
177/177 [=====] - 0s 1ms/step - loss: 0.3655 -
accuracy: 0.8273

Epoch 86/100
177/177 [=====] - 0s 1ms/step - loss: 0.3653 -
accuracy: 0.8259

Epoch 87/100
177/177 [=====] - 0s 1ms/step - loss: 0.3651 -
accuracy: 0.8278

Epoch 88/100
177/177 [=====] - 0s 1ms/step - loss: 0.3653 -
accuracy: 0.8241

Epoch 89/100
177/177 [=====] - 0s 1ms/step - loss: 0.3637 -
accuracy: 0.8289

Epoch 90/100
177/177 [=====] - 0s 1ms/step - loss: 0.3647 -
accuracy: 0.8284

Epoch 91/100
177/177 [=====] - 0s 1ms/step - loss: 0.3630 -
accuracy: 0.8294

Epoch 92/100
177/177 [=====] - 0s 1ms/step - loss: 0.3647 -
accuracy: 0.8271

Epoch 93/100
177/177 [=====] - 0s 1ms/step - loss: 0.3635 -
accuracy: 0.8246

Epoch 94/100
177/177 [=====] - 0s 1ms/step - loss: 0.3623 -
accuracy: 0.8269

Epoch 95/100
177/177 [=====] - 0s 1ms/step - loss: 0.3625 -
accuracy: 0.8271

```
Epoch 96/100
177/177 [=====] - 0s 1ms/step - loss: 0.3611 -
accuracy: 0.8277
Epoch 97/100
177/177 [=====] - 0s 1ms/step - loss: 0.3608 -
accuracy: 0.8291
Epoch 98/100
177/177 [=====] - 0s 1ms/step - loss: 0.3611 -
accuracy: 0.8300
Epoch 99/100
177/177 [=====] - 0s 1ms/step - loss: 0.3616 -
accuracy: 0.8280
Epoch 100/100
177/177 [=====] - 0s 1ms/step - loss: 0.3607 -
accuracy: 0.8285
```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x7f877bfcde80>
```

```
[ ]: model.evaluate(X_test, y_test)
```

```
45/45 [=====] - 0s 949us/step - loss: 0.4588 -
accuracy: 0.7779
```

```
[ ]: [0.458802729845047, 0.7778566479682922]
```

```
[ ]: yp = model.predict(X_test)
yp[:5]
```

```
[ ]: array([[0.36123043],
          [0.50069755],
          [0.32254955],
          [0.76979315],
          [0.09441373]], dtype=float32)
```

```
[ ]: y_pred = []
for element in yp:
    if element > 0.5:
        y_pred.append(1)
    else:
        y_pred.append(0)
```

```
[ ]: y_pred[:10]
```

```
[ ]: [0, 1, 0, 1, 0, 1, 1, 1, 0, 0]
```

```
[ ]: y_test[:10]
```



```
[ ]: 4213    1
      5035    0
      3713    1
      1720    0
      234     0
      4558    1
      40      0
      3455    1
      5944    1
      1089    0
      Name: Churn, dtype: int64
```

Confusion matrix

```
[ ]: from sklearn.metrics import confusion_matrix , classification_report


      print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.86	0.85	1023
1	0.60	0.57	0.59	386
accuracy			0.78	1409
macro avg	0.72	0.71	0.72	1409
weighted avg	0.77	0.78	0.78	1409

```
[ ]: cm = tf.math.confusion_matrix(labels=y_test,predictions=y_pred)

      plt.figure(figsize = (10,7))
      sns.heatmap(cm, annot=True, fmt='d')
      plt.xlabel('Predicted')
      plt.ylabel('Truth')
```

```
[ ]: Text(69.0, 0.5, 'Truth')
```



```
[ ]: y_test.shape
```

```
[ ]: (1409,)
```

Accuracy

```
[ ]: round((862+229)/(862+229+137+179),2)
```

```
[ ]: 0.78
```

Precision for 0 class. i.e. Precision for customers who did not churn

```
[ ]: round(862/(862+179),2)
```

```
[ ]: 0.83
```

Precision for 1 class. i.e. Precision for customers who actually churned

```
[ ]: round(229/(229+137),2)
```

```
[ ]: 0.63
```

Recall for 0 class

```
[ ]: round(862/(862+137),2)
```

```
[ ]: 0.86
```

```
[ ]: round(229/(229+179),2)
```

```
[ ]: 0.56
```

```
[ ]:
```